

= ChatGPT ▾

✖ Memory full ⌂ ...

I have AWS SDK professional training for 20 years laterals

Since your audience has **20+ years of experience**, the training must be **architecture-focused, enterprise-level, and practical**, not beginner-level SDK syntax. Senior professionals care about **design patterns, scalability, automation, performance, and real-world enterprise integration**.

Here is a **complete AWS SDK Professional Training Plan (Senior / 20+ Years Experience)**:

AWS SDK Professional Training Plan (For Senior Professionals)

Training Objective

Enable experienced professionals to:

- Architect enterprise-grade cloud applications using AWS SDK
 - Automate infrastructure and workflows programmatically
 - Build scalable, secure, production-ready systems
 - Integrate AWS services into enterprise applications
-

Module 1: AWS SDK Architecture Deep Dive

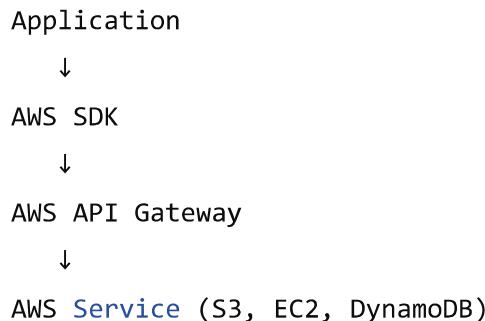
Topics

- AWS SDK internal architecture
- SDK vs CLI vs Console vs IaC
- When to use SDK vs Terraform vs CloudFormation
- SDK lifecycle and request flow

Architecture Flow

java

Copy code



Demo (Python boto3)

python

 Copy code

```
import boto3

s3 = boto3.client('s3')
response = s3.list_buckets()

for bucket in response['Buckets']:
    print(bucket['Name'])
```

Enterprise Use Case

ERP system listing storage buckets dynamically.

Module 2: Enterprise Authentication and Security

Topics

- IAM Roles vs IAM Users
- AssumeRole pattern
- Temporary credentials (STS)
- Cross-account access

Demo: Assume Role

python

 Copy code

```
import boto3
```

```
sts = boto3.client('sts')

response = sts.assume_role(
    RoleArn="arn:aws:iam::123456789012:role/AdminRole",
    RoleSessionName="Session1"
)

print(response['Credentials'])
```

Enterprise Use Case

Finance system accessing secure AWS account.

Module 3: SDK with Compute Services

Topics

- Launch EC2 instances via SDK
- Auto Scaling automation
- Lambda invocation via SDK

Demo: Launch EC2

python

 Copy code

```
ec2 = boto3.resource('ec2')

instance = ec2.create_instances(
    ImageId='ami-12345',
    MinCount=1,
    MaxCount=1,
    InstanceType='t2.micro'
)

print(instance[0].id)
```

Enterprise Use Case

Auto-provision servers during peak load.

Module 4: SDK with Storage Services

Topics

- Upload/download files programmatically
- Backup automation
- Multi-region replication

Demo: Upload file

python

 Copy code

```
s3 = boto3.client('s3')

s3.upload_file(
    'report.pdf',
    'company-bucket',
    'reports/report.pdf'
)
```

Enterprise Use Case

Automated backup system.

Module 5: SDK with Database Services

Topics

- DynamoDB CRUD via SDK
- RDS management
- Query optimization

Demo: Insert into DynamoDB

python

 Copy code

```
dynamodb = boto3.resource('dynamodb')
table = dynamodb.Table('Employees')

table.put_item(
    Item={
```

```
'ID': '101',
'Name': 'Robert',
'Department': 'Finance'
}
)
```

Module 6: Event-Driven Enterprise Automation

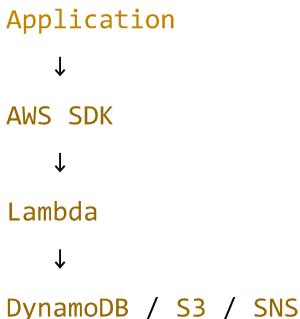
Topics

- SDK with Lambda
- SDK with EventBridge
- Automation workflows

Architecture

mathematica

 Copy code



Module 7: Enterprise Integration Patterns

Topics

- SDK in Microservices
- SDK in enterprise backend systems
- SDK with message queues (SQS)

Demo: Send message to queue

python

 Copy code

```
    sqs = boto3.client('sns')

    sqs.send_message(
        QueueUrl='queue-url',
        MessageBody='Order created'
    )
```

Module 8: Error Handling and Resilience

Topics

- Retry strategies
- Exponential backoff
- Fault tolerance

Demo

python

 Copy code

```
import botocore

try:
    s3.list_buckets()
except botocore.exceptions.ClientError as e:
    print("Error:", e)
```

Module 9: Performance Optimization

Topics

- Connection pooling
- Parallel processing
- Async SDK usage

Module 10: Multi-Account Enterprise Architecture

Topics

- Cross-account automation
 - Enterprise governance
 - Role chaining
-

Module 11: SDK in DevOps and Automation

Topics

- CI/CD integration
 - Automated deployments
 - Infrastructure provisioning
-

Module 12: Real-World Enterprise Case Study

Case Study: Automated Enterprise Backup System

Flow:

mathematica

 Copy code

Application

↓

SDK

↓

Lambda

↓

S3 Backup

↓

Notification via SNS

Advanced Enterprise Demo (End-to-End)

Automate:

- Launch EC2

- Upload logs to S3
 - Store metadata in DynamoDB
-

Enterprise-Level Capstone Project

Project: Enterprise Infrastructure Automation System

Features:

- Auto-create EC2 servers
 - Store metadata in DynamoDB
 - Upload logs to S3
 - Send notifications via SNS
-

Recommended SDK Languages for Enterprise

Priority order:

1. Python (boto3)
 2. Java SDK
 3. Node.js SDK
 4. .NET SDK
-

Training Duration Plan (Recommended)

Module	Duration
Architecture	2 hours
Security	2 hours
Compute	3 hours
Storage	2 hours
Database	2 hours

Module	Duration
Event-Driven	2 hours
Integration	2 hours
Resilience	1.5 hours
Performance	1.5 hours
DevOps	2 hours
Capstone	4 hours

Total: 24 Hours

What Senior Professionals Expect (Important)

Focus on:

- Architecture diagrams
- Real enterprise use cases
- Automation patterns
- Performance and scaling
- Security best practices
- Failure scenarios

NOT basic syntax.

If you want, I can also create a 5-day AWS SDK enterprise training plan with demos, architecture diagrams, and hands-on labs specifically tailored for 20-year experienced professionals.



I have DAY 1 -> TOC below -> I wanted professional Demos and Lab steps -> also examples, scenarios: Day 1 – Introduction to AWS Cloud & Developer Role

- AWS Overview & Free Tier Setup
 - IAM Basics (Users, Groups, Policies)
 - Developer Tools: AWS Console, CLI, SDKs
- Hands-on: Free Tier setup + IAM user creation + AWS CLI configuration

Below is a **Day 1 Professional Training Plan with Enterprise-Level Demos, Labs, Scenarios, and Architecture Thinking**, specifically designed for **20+ years experienced professionals**.

Day 1 – Introduction to AWS Cloud & Developer Role (Professional Level)

Learning Outcome (Senior-Level)

By end of Day 1, participants will:

- Understand AWS architecture from an enterprise developer perspective
- Configure secure developer access using IAM best practices
- Use AWS CLI and SDK professionally (not console-only usage)
- Implement enterprise-grade authentication model

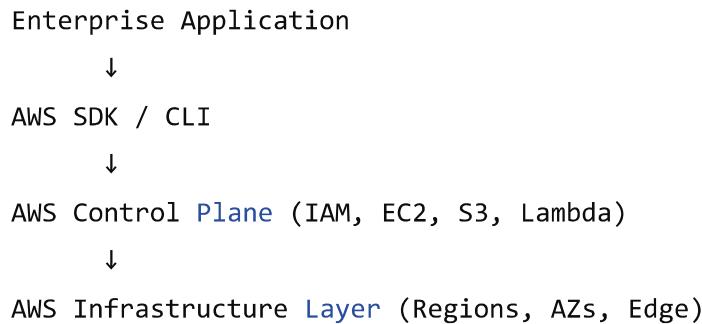
Part 1: AWS Overview & Free Tier Setup (Professional Perspective)

Concept Focus (Not basic)

Explain AWS in terms of enterprise architecture layers:

java

 Copy code

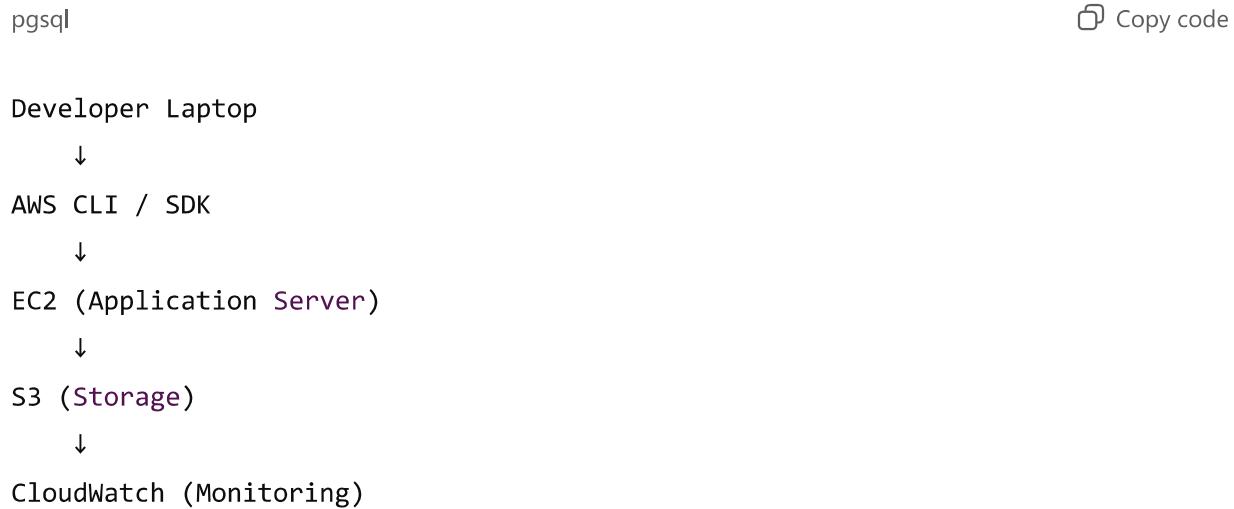


Professional Demo 1: Enterprise Application Hosting Architecture

Scenario

Your company is migrating an on-prem application to AWS.

Architecture:



Explain:

Developers interact using:

- Console → Manual
- CLI → Automation
- SDK → Application integration

Enterprise uses mostly CLI and SDK.

Lab 1: AWS Free Tier Setup (Professional Way)

Step 1: Create AWS Account

Go to:

<https://aws.amazon.com/free> ↗

Configure:

- Root user email
 - Enable MFA (mandatory enterprise practice)
-

Step 2: Setup Billing Alert (Critical Enterprise Step)

Go to:

Billing → Budgets → Create Budget

Set:

makefile

 Copy code

Budget: \$5

Alert: Email notification

Enterprise reason:

Prevent unexpected cloud costs.

Part 2: IAM Basics (Enterprise Security Model)

Explain IAM from enterprise architecture perspective:

Never use Root user.

Use model:

pgsql

 Copy code

```
Root User (Locked)
```

↓

```
Admin User
```

↓

```
Developer Users
```

↓

```
Roles (Recommended)
```

Professional Demo 2: Enterprise IAM Model

Scenario

Company has:

- 100 developers
- Each needs controlled access

Use IAM Users + Groups.

Architecture:

```
pgsql
```

 Copy code

```
IAM Group: Developers
```

↓

```
IAM Policy: EC2 + S3 Access
```

↓

```
IAM Users assigned to group
```

Lab 2: Create Enterprise IAM Structure

Step 1: Create IAM Group

Go to:

IAM → User Groups → Create group

Name:

nginx

 Copy code

Developers

Attach Policy:

nginx

 Copy code

AmazonEC2FullAccess

AmazonS3FullAccess

Step 2: Create IAM User

IAM → Users → Create user

Name:

dev-user1

 Copy code

Select:

pgsql

 Copy code

Programmatic Access

AWS Console Access

Add to group:

nginx

 Copy code

Developers

Download credentials.

Enterprise Note:

Never share credentials via email. Use secure vault.

Professional Demo 3: Access Validation Test

Login using IAM user.

Verify:

Can access EC2 and S3

Cannot access Billing (important security validation)

Part 3: Developer Tools: Console vs CLI vs SDK

Explain enterprise usage comparison:

Tool	Enterprise Usage
Console	Manual tasks, debugging
CLI	Automation, DevOps
SDK	Application integration

Professional Demo 4: AWS CLI Installation and Setup

Step 1: Install AWS CLI

Windows:

Download:

<https://aws.amazon.com/cli/> ↗

Verify:

css

 Copy code

```
aws --version
```

Expected:

 Copy code

```
aws-cli/2.x.x
```

Step 2: Configure CLI

Run:

```
nginx
```

 Copy code

```
aws configure
```

Enter:

```
vbnf
```

 Copy code

```
Access Key
```

```
Secret Key
```

```
Region: ap-south-1
```

```
Output: json
```

Enterprise Explanation:

CLI stores credentials securely locally.

Lab 3: Professional CLI Hands-on

Test 1: List S3 Buckets

```
bash
```

 Copy code

```
aws s3 ls
```

Expected output:

```
php
```

 Copy code

```
(empty or bucket list)
```

Test 2: Create Bucket

```
pgsql|
```

 Copy code

```
aws s3 mb s3://company-dev-bucket-unique-name
```

Verify:

```
bash|
```

 Copy code

```
aws s3 ls
```

Test 3: Upload File

Create file:

```
test.txt|
```

 Copy code

Upload:

```
bash|
```

 Copy code

```
aws s3 cp test.txt s3://company-dev-bucket-unique-name
```

Enterprise Use Case:

Application logs uploaded automatically.

Professional Demo 5: EC2 Creation via CLI

Run:

```
sql|
```

 Copy code

```
aws ec2 describe-instances
```

Explain:

Developers automate infrastructure.

Part 4: SDK Introduction (Professional Developer Perspective)

Explain:

SDK used inside application code.

Architecture:

markdown

 Copy code

Java/Python Application

↓

AWS SDK

↓

AWS Services

Professional Demo 6: SDK Example (Python)

Install:

nginx

 Copy code

`pip install boto3`

Code:

python

 Copy code

```
import boto3
```

```
s3 = boto3.client('s3')
```

```
response = s3.list_buckets()
```

```
for bucket in response['Buckets']:
    print(bucket['Name'])
```

Enterprise Scenario:

Application dynamically reads storage.

Professional Scenario 1: Enterprise Developer Workflow

Developer builds application that:

- Uploads reports to S3
- Launches EC2 automatically
- Stores data in database

Workflow:

mathematica

 Copy code

Application



AWS SDK



AWS Infrastructure

Professional Scenario 2: Secure Developer Access Model

Enterprise Security Architecture:

pgsql

 Copy code

Developer



IAM User



IAM Role



AWS Resources

Not direct root access.

Lab 4: Real Enterprise Lab (End-to-End)

Objective:

Simulate enterprise developer setup.

Step 1

Create IAM User

Step 2

Configure CLI

Step 3

Create S3 bucket via CLI

Step 4

Upload file via CLI

Step 5

Access via SDK code

Expected Outcome:

Fully functional enterprise developer environment.

Professional Troubleshooting Demo

Common errors and fixes:

Error:

nginx

 Copy code

AccessDenied

Cause:

IAM policy missing.

Fix:

Attach correct policy.

Error:

nginx

 Copy code

InvalidAccessKeyId

Fix:

Reconfigure CLI.

Enterprise Best Practices (Important for Senior Professionals)

Always use:

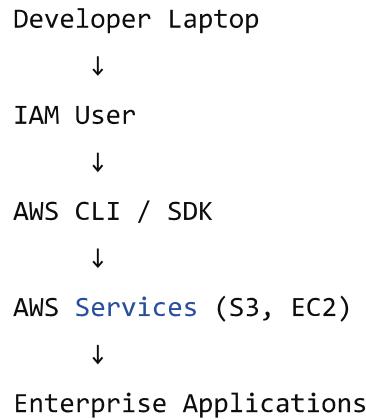
- IAM Users or Roles (Never root)
- MFA enabled

- Least privilege access
 - CLI or SDK for automation
 - Billing alerts enabled
-

End of Day 1 Architecture Summary

java

 Copy code



Day 1 Professional Capstone Lab

Participants will:

- Create IAM User
 - Configure CLI
 - Create S3 bucket
 - Upload file
 - Access using SDK
-

If you want, I can also prepare Day 1 PowerPoint slides, lab guide PDF, and trainer demo script specifically tailored for senior professionals.



I have DAY 2 -> TOC below -> I wanted professional Demos and Lab steps -> also examples, scenarios: Day 2 – Terraform Infrastructure Automation and ECS Setup ()

- Infrastructure as Code Concepts
- Manual console setup vs declarative provisioning.
- Version-controlled, repeatable deployments.
- Terraform Workflow Overview
- init → plan → apply → destroy cycle.
- What Terraform means by “desired state.”
- Terraform Files & Structure
- Terraform Variables (Theory)
- Purpose of variables (reusability & customization).
- Types: string, number, bool, map, list.
- Variable precedence (defaults > tfvars > CLI > env).
- Create the AWS Resource using KMS using Terraform Code

Below is a **Day 2 – Professional-Level Training Plan** designed for senior/lateral professionals with strong architecture thinking.

This is not basic Terraform syntax — this is **enterprise IaC mindset + production scenarios + hands-on labs.**

Day 2 – Terraform Infrastructure Automation & ECS Setup

Learning Outcome (Senior-Level)

By end of Day 2, participants will:

- Understand Infrastructure as Code from enterprise architecture perspective
- Compare manual vs declarative provisioning (governance impact)
- Master Terraform workflow (init → plan → apply → destroy)
- Understand desired state model deeply
- Implement reusable Terraform structure
- Use variables professionally

- Create AWS KMS key via Terraform
 - Provision ECS cluster using Terraform
-

PART 1 – Infrastructure as Code (Enterprise Thinking)

Professional Explanation

Infrastructure as Code (IaC) means:

Infrastructure defined as version-controlled code.

Traditional model:

arduino

 Copy code

Engineer → **Console** → Click → Resource created

Enterprise Problem:

- No tracking
- No reproducibility
- Human error
- No audit

Terraform model:

css

 Copy code

Git Repo → Terraform **Code** → Plan → Apply → Infrastructure

Enterprise Advantage:

- Version control
 - Peer review
 - Repeatable environments
 - Disaster recovery ready
-

Professional Demo 1: Manual Console vs Terraform

Scenario

Company needs 3 environments:

- Dev
- QA
- Prod

Manual Console Issues:

- Different configurations
- Missed tags
- Wrong security group

Terraform Solution:

Single codebase with variable-based environments.

PART 2 – Terraform Workflow Deep Dive

Terraform Lifecycle

csharp

 Copy code

```
terraform init  
terraform plan  
terraform apply  
terraform destroy
```

Professional Demo 2 – Basic Terraform Workflow

Step 1: Install Terraform

Verify:

nginx

 Copy code

```
terraform -version
```

Step 2: Create Project Structure

css

 Copy code

```
aws-terraform-project/
├── main.tf
├── variables.tf
├── terraform.tfvars
└── outputs.tf
```

Lab 1 – First Terraform Resource (S3 Example)

main.tf

hcl

 Copy code

```
provider "aws" {
    region = "ap-south-1"
}

resource "aws_s3_bucket" "company_bucket" {
    bucket = "company-dev-bucket-unique-12345"

    tags = {
        Environment = "Dev"
        Owner       = "PlatformTeam"
    }
}
```

Step 3: Run Workflow

Initialize

csharp

 Copy code

```
terraform init
```

See Execution Plan

nginx

 Copy code

```
terraform plan
```

Explain:

Terraform compares:

- Current state
- Desired state (code)

Shows delta.

Apply

nginx

 Copy code

```
terraform apply
```

Confirm: yes

Destroy

nginx

 Copy code

```
terraform destroy
```

Enterprise Use Case:

Environment teardown after testing.

PART 3 – Desired State Model (Deep Understanding)

Terraform does NOT say:

"Create bucket"

It says:

"This is what infrastructure should look like."

Terraform then:

- Compares real world
- Makes changes to match desired state

Enterprise Example:

If bucket deleted manually,

Next `terraform apply` → Recreates bucket.

This ensures compliance.

PART 4 – Terraform File Structure (Professional Standard)

Enterprise Structure:

css

 Copy code

```
project/
  ├── main.tf
  ├── variables.tf
  ├── outputs.tf
  ├── terraform.tfvars
  ├── backend.tf
  └── modules/
    ├── ecs/
    ├── kms/
    └── vpc/
```

Why modular?

- Reusability
- Team collaboration
- Governance

PART 5 – Terraform Variables (Professional Usage)

Why Variables?

Avoid hardcoding:

Bad:

ini

 Copy code

```
region = "ap-south-1"
```

Good:

ini

 Copy code

```
region = var.aws_region
```

variables.tf

hcl

 Copy code

```
variable "aws_region" {  
    type      = string  
    description = "AWS Region"  
}
```

```
variable "environment" {  
    type      = string  
    default  = "dev"  
}
```

```
variable "enable_deletion" {  
    type      = bool  
    default  = false  
}
```

terraform.tfvars

hcl

Copy code

```
aws_region = "ap-south-1"  
environment = "dev"  
enable_deletion = true
```

Variable Precedence (Enterprise Exam Concept)

Lowest → Highest:

1. Default in variable block
2. terraform.tfvars
3. CLI -var
4. Environment variables

Enterprise Scenario:

CI/CD pipeline overrides region via CLI.

PART 6 – Create AWS KMS Key Using Terraform

Enterprise Scenario

Finance application needs encryption.

Requirement:

- Centralized KMS key
 - Rotation enabled
 - Tagged properly
-

Lab 2 – Create KMS Key

main.tf

hcl

Copy code

```

provider "aws" {
    region = var.aws_region
}

resource "aws_kms_key" "company_kms" {
    description          = "Enterprise KMS Key"
    deletion_window_in_days = 7
    enable_key_rotation     = true

    tags = {
        Environment = var.environment
        Owner       = "SecurityTeam"
    }
}

```

outputs.tf

hcl

 Copy code

```

output "kms_key_id" {
    value = aws_kms_key.company_kms.id
}

```

Run:

csharp

 Copy code

```

terraform init
terraform plan
terraform apply

```

Verify in AWS Console → KMS → Keys

Enterprise Discussion:

Why enable rotation?

Compliance + audit requirements.

PART 7 – ECS Setup Using Terraform

Enterprise Scenario

Company wants containerized microservice platform.

Architecture:

pgsql

 Copy code

ECS Cluster

↓

Task Definition

↓

Service

↓

Load Balancer

Lab 3 – Create ECS Cluster

main.tf

hcl

 Copy code

```
resource "aws_ecs_cluster" "company_cluster" {  
    name = "company-ecs-cluster"  
}
```

Apply:

nginx

 Copy code

```
terraform apply
```

Verify in ECS Console.

Professional Scenario – Real Enterprise Deployment

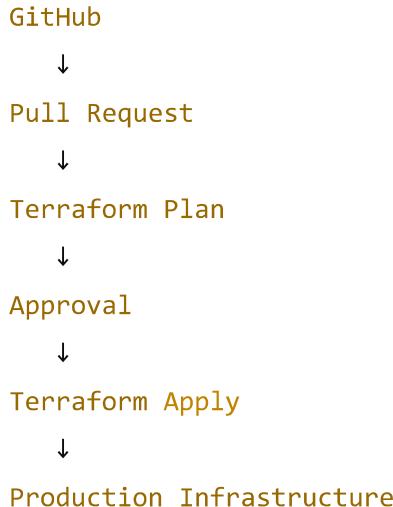
Company runs:

- 10 microservices
- 3 environments
- Central encryption via KMS
- ECS-based architecture

Terraform Flow:

mathematica

 Copy code



PART 8 – End-to-End Professional Lab

Objective:

Provision:

- KMS key
- ECS cluster
- S3 bucket (logs)

Lab Steps Summary

1. Create folder structure
2. Create variables.tf
3. Create main.tf

4. Add KMS resource
5. Add ECS cluster resource
6. Run:

csharp

 Copy code

```
terraform init  
terraform plan  
terraform apply
```

Troubleshooting Demo (Important for Seniors)

Error:

pgsql

 Copy code

```
Error: No valid credential sources found
```

Fix:

Configure AWS CLI first.

Error:

nginx

 Copy code

```
BucketAlreadyExists
```

Fix:

Use unique name.

Enterprise Best Practices (Very Important)

- Use remote backend (S3 + DynamoDB)
- Use state locking
- Never commit secrets
- Use modules

- Tag everything
 - Use separate workspaces per environment
-

Architecture Summary – Day 2

CSS

 Copy code

Terraform [Code](#)



Terraform State



AWS API



Infrastructure (KMS, ECS, S3)

Day 2 Professional Capstone

Build:

- Secure container platform
 - Encrypted infrastructure
 - Version-controlled IaC
-

If you want, I can prepare **Day 3 – ECS Full Deployment with Task Definition, Service, ALB, IAM Roles, and CI/CD Integration** at enterprise level.



I have DAY 3 -> TOC below -> I wanted professional Demos and Lab steps -> also examples, scenarios: Day 3 – AWS CLI & SDK (Node.js or Python)

- CLI Scripting for EC2, RDS, IAM
- SDK Authentication and API Calls

Hands-on: Script EC2/RDS creation and test SDK query

- Why use the AWS CLI?
 - Repeatable automation, faster testing, integration with shell scripts.
- CLI Configuration Concepts
 - Access Key ID, Secret Key, Region, Output Format.
 - Profiles (default vs custom).
- Command Structure Pattern– How to find help and documentation.
- CLI vs Console vs SDK
 - When to choose each for automation.
- Security Considerations
 - IAM least-privilege roles, avoiding hard-coded keys, using SSO and environment variables.

Below is a Day 3 – Professional-Level Training Plan (AWS CLI & SDK) designed specifically for 20+ years experienced professionals, with enterprise demos, automation labs, real-world scenarios, and production practices.

This focuses on automation, scripting, secure authentication, and programmatic infrastructure control.

Day 3 – AWS CLI & SDK Automation (Enterprise-Level)

Learning Outcome

By end of Day 3, participants will:

- Automate infrastructure using AWS CLI scripts
 - Understand CLI command structure deeply
 - Use profiles and secure authentication methods
 - Automate EC2, RDS, IAM using CLI
 - Integrate SDK into enterprise applications
 - Implement secure SDK authentication
 - Understand CLI vs SDK vs Console usage in production
-

PART 1 – Why Enterprise Uses AWS CLI

Enterprise Scenario

Developer needs to create:

- 10 EC2 servers
- Multiple environments (Dev, QA, Prod)

Manual Console → Slow and error-prone

CLI → Automated and repeatable

Enterprise Automation Flow:

markdown

 Copy code

Shell Script

↓

AWS CLI

↓

AWS API

↓

Infrastructure Created

Professional Demo 1 – AWS CLI Configuration (Enterprise Way)

Step 1: Configure CLI

Run:

bash

 Copy code

aws configure

Enter:

vbnet

 Copy code

Access **Key** ID

Secret Access **Key**

Region: ap-south-1

Output: json

Step 2: Verify Access

bash

 Copy code

```
aws sts get-caller-identity
```

Expected Output:

json

 Copy code

```
{  
    "UserId": "...",  
    "Account": "...",  
    "Arn": "arn:aws:iam::xxxx:user/dev-user"  
}
```

Enterprise Use Case:

Verify identity before running automation scripts.

PART 2 – CLI Profiles (Enterprise Multi-Account Setup)

Enterprise Scenario:

Company has:

- Dev Account
- QA Account
- Prod Account

Create profile:

bash

 Copy code

```
aws configure --profile dev
```

Use profile:

bash

 Copy code

```
aws s3 ls --profile dev
```

Enterprise Advantage:

Supports multi-account automation.

PART 3 – CLI Command Structure Pattern

Standard Structure:

bash

 Copy code

```
aws <service> <operation> <parameters>
```

Examples:

List EC2 instances:

bash

 Copy code

```
aws ec2 describe-instances
```

List IAM users:

bash

 Copy code

```
aws iam list-users
```

Professional Demo 2 – Finding Help

Command help:

bash

 Copy code

```
aws ec2 help
```

Specific help:

bash

 Copy code

```
aws ec2 run-instances help
```

Enterprise Use Case:

Developers explore APIs quickly.

PART 4 – Automate EC2 Creation Using CLI

Enterprise Scenario

Auto-create server for application deployment.

Lab 1 – Create EC2 via CLI

Step 1: Create Security Group

bash

 Copy code

```
aws ec2 create-security-group \
--group-name dev-sg \
--description "Dev Security Group"
```

Step 2: Launch EC2 Instance

bash

 Copy code

```
aws ec2 run-instances \
--image-id ami-0f58b397bc5c1f2e8 \
--count 1 \
--instance-type t2.micro
```

Step 3: Verify Instance

bash

 Copy code

```
aws ec2 describe-instances
```

Enterprise Automation Scenario:

DevOps pipeline auto-creates servers.

PART 5 – Automate IAM Using CLI

Lab 2 – Create IAM User via CLI

bash

 Copy code

```
aws iam create-user --user-name automation-user
```

Attach policy:

bash

 Copy code

```
aws iam attach-user-policy \  
--user-name automation-user \  
--policy-arn arn:aws:iam::aws:policy/AmazonEC2FullAccess
```

Enterprise Scenario:

Automated onboarding system.

PART 6 – Automate RDS Creation Using CLI

Enterprise Scenario

Application requires database automatically during deployment.

Lab 3 – Create RDS Instance

bash

Copy code

```
aws rds create-db-instance \
--db-instance-identifier dev-db \
--db-instance-class db.t3.micro \
--engine mysql \
--master-username admin \
--master-user-password Password123 \
--allocated-storage 20
```

Enterprise Use Case:

Automated database provisioning.

PART 7 – CLI Automation Script (Professional Demo)

Create file:

deploy.sh

bash

 Copy code

```
#!/bin/bash

echo "Creating EC2 instance..."
aws ec2 run-instances \
--image-id ami-0f58b397bc5c1f2e8 \
--count 1 \
--instance-type t2.micro

echo "Listing instances..."
aws ec2 describe-instances
```

Run:

bash

 Copy code

```
chmod +x deploy.sh
./deploy.sh
```

Enterprise Use Case:
CI/CD automation.

PART 8 – SDK Authentication and API Calls

SDK is used inside applications.

Enterprise Architecture:

markdown

 Copy code

Enterprise Application

↓

AWS SDK

↓

AWS API

Professional Demo 3 – SDK Setup (Python)

Install:

bash

 Copy code

pip install boto3

Lab 4 – SDK Query EC2 Instances

Python code:

python

 Copy code

```
import boto3
```

```
ec2 = boto3.client('ec2')
```

```
response = ec2.describe_instances()
```

```
for reservation in response['Reservations']:
    for instance in reservation['Instances']:
        print(instance['InstanceId'])
```

Enterprise Use Case:

Application monitors infrastructure.

Professional Demo 4 – Create EC2 Using SDK

python

 Copy code

```
import boto3

ec2 = boto3.resource('ec2')

instance = ec2.create_instances(
    ImageId='ami-0f58b397bc5c1f2e8',
    MinCount=1,
    MaxCount=1,
    InstanceType='t2.micro'
)

print("Instance created:", instance[0].id)
```

PART 9 – CLI vs Console vs SDK (Enterprise Comparison)

Tool	Use Case
Console	Manual troubleshooting
CLI	Automation scripts
SDK	Application integration

Enterprise Flow:

arduino

Copy code

Application → SDK → AWS

DevOps → CLI → AWS

Admin → **Console** → AWS

PART 10 – Security Best Practices (Critical Enterprise Topic)

Never hardcode keys like:

python

 Copy code

```
aws_access_key_id="ABC"
```

Use:

CLI configured credentials

OR environment variables:

bash

 Copy code

```
export AWS_ACCESS_KEY_ID=xxx  
export AWS_SECRET_ACCESS_KEY=xxx
```

OR IAM roles (best practice)

Professional Demo 5 – Using Environment Variables

Linux:

bash

 Copy code

```
export AWS_ACCESS_KEY_ID=your_key  
export AWS_SECRET_ACCESS_KEY=your_secret
```

Verify:

bash

 Copy code

```
aws sts get-caller-identity
```

Enterprise Scenario – Production Application

Application running on EC2 uses IAM role:

```
pgsql
```

 Copy code

EC2 Instance

↓

IAM Role

↓

AWS SDK

↓

S3 Access

No keys stored.

Secure architecture.

End-to-End Professional Lab

Participants will:

Step 1:

Configure CLI

Step 2:

Create EC2 using CLI

Step 3:

Create IAM user using CLI

Step 4:

Create RDS using CLI

Step 5:

Run SDK script to query EC2

Enterprise Architecture Summary

markdown

 Copy code

Shell Script

↓

AWS CLI

↓

AWS API

↓

Infrastructure

Application

↓

AWS SDK

↓

AWS API

Professional Troubleshooting Demo

Error:

css

 Copy code

Unable **to** locate credentials

Fix:

Run:

bash

 Copy code

aws configure

Error:

nginx

 Copy code

AccessDenied

Fix:

Attach proper IAM policy.

Professional Capstone Lab – Enterprise Automation

Create script that:

- Creates EC2
- Creates IAM user
- Queries using SDK

Fully automated deployment.

If you want, I can also create **real enterprise-level CI/CD automation using CLI + SDK + Terraform integrated together**, which is what senior professionals expect in real companies.



I have DAY 3 -> TOC below -> I wanted professional Demos and Lab steps -> also examples, scenarios: Day 4 – Amazon S3 Deep Dive

Introduction to Object Storage

- Difference between block, file, and object storage.
- Why S3 is a backbone for data storage in AWS.

Key S3 Concepts

- Buckets, Objects, and Regions.
- Object structure (key, value, metadata, version ID).

Storage Classes & Use Cases

- Standard, Intelligent-Tiering, Glacier, Deep Archive.

- Real-time example: Storing app logs vs archived backups.
- S3 Features Overview
- Versioning, Lifecycle Management, Encryption (SSE-S3, SSE-KMS).
 - Access Control: Bucket Policy, ACL, IAM Roles.
- Cost Optimization Basics
- Data transfer cost, storage class transitions.

Below is a Day 4 – Amazon S3 Deep Dive (Professional-Level Training) with enterprise demos, architecture scenarios, hands-on labs, automation examples, and real-world best practices, designed specifically for senior professionals (20+ years experience).

This focuses on enterprise storage architecture, security, lifecycle automation, and cost optimization.

Day 4 – Amazon S3 Deep Dive (Enterprise-Level)

Learning Outcomes

By end of Day 4, participants will:

- Understand storage architecture types (block, file, object)
 - Design enterprise storage using S3
 - Implement versioning, encryption, lifecycle policies
 - Secure S3 using IAM roles and bucket policies
 - Optimize storage costs using storage classes
 - Automate S3 using CLI and SDK
-

PART 1 – Storage Types Architecture

Enterprise Storage Comparison

Storage Type	Example	AWS Service	Use Case
Block Storage	Hard disk	EBS	Databases

Storage Type	Example	AWS Service	Use Case
File Storage	Network drive	EFS	Shared folders
Object Storage	Cloud storage	S3	Logs, backups, media

Professional Architecture Example

Enterprise Application Storage Model:

mathematica

 Copy code

Application Server (EC2)



EBS → Database storage

Application



EFS → Shared files

Application



S3 → Logs, backups, media files

Professional Demo 1 – Why S3 is Backbone

Enterprise systems use S3 for:

- Application logs
- Backups
- Data lakes
- Machine learning datasets
- Static websites

Example Architecture:

mathematica

 Copy code



PART 2 – S3 Key Concepts

Core Components

pgsql

Copy code

Bucket → Container

Object → File

Key → File path/name

Metadata → File info

Version ID → File version

Example Object Structure:

vbnet

Copy code

Bucket: company-data

Object Key:

logs/2026/app.log

Metadata:

Created date

Owner

Version ID

Lab 1 – Create S3 Bucket (Enterprise Standard)

Using CLI

bash

 Copy code

```
aws s3 mb s3://company-prod-storage-unique123
```

Verify:

bash

 Copy code

```
aws s3 ls
```

Upload Object

bash

 Copy code

```
aws s3 cp test.txt s3://company-prod-storage-unique123/logs/test.txt
```

Verify:

bash

 Copy code

```
aws s3 ls s3://company-prod-storage-unique123/logs/
```

Enterprise Scenario

Application uploads logs automatically.

PART 3 – Storage Classes and Enterprise Use Cases

Storage Class Comparison

Storage Class	Use Case
Standard	Active application data
Intelligent-Tiering	Unknown access patterns

Storage Class	Use Case
Glacier	Backup storage
Deep Archive	Long-term archive

Enterprise Example

Application logs lifecycle:

sql

 Copy code

```
Day 0-30 → Standard  
Day 30-90 → Intelligent-Tiering  
Day 90-365 → Glacier  
After 365 → Deep Archive
```

Cost reduction up to 90%.

Professional Demo 2 – Upload with Storage Class

bash

 Copy code

```
aws s3 cp test.txt s3://company-prod-storage-unique123 \  
--storage-class STANDARD_IA
```

PART 4 – Versioning (Critical Enterprise Feature)

Why Versioning?

Protects against:

- Accidental deletion
- Overwrites
- Ransomware

Lab 2 – Enable Versioning

bash

 Copy code

```
aws s3api put-bucket-versioning \
--bucket company-prod-storage-unique123 \
--versioning-configuration Status=Enabled
```

Test Versioning

Upload file:

bash

 Copy code

```
aws s3 cp test.txt s3://company-prod-storage-unique123
```

Modify file and upload again.

List versions:

bash

 Copy code

```
aws s3api list-object-versions \
--bucket company-prod-storage-unique123
```

Enterprise Benefit:

Recover any version.

PART 5 – Lifecycle Management (Cost Optimization)

Enterprise Scenario:

Move old files automatically to Glacier.

Lab 3 – Lifecycle Policy

Create file lifecycle.json

json

 Copy code

```
{  
  "Rules": [  
    {  
      "ID": "ArchiveRule",  
      "Prefix": "",  
      "Status": "Enabled",  
      "Transitions": [  
        {  
          "Days": 30,  
          "StorageClass": "GLACIER"  
        }  
      ]  
    }  
  ]  
}
```

Apply policy:

bash

 Copy code

```
aws s3api put-bucket-lifecycle-configuration \  
--bucket company-prod-storage-unique123 \  
--lifecycle-configuration file://lifecycle.json
```

Enterprise Benefit:

Fully automated cost optimization.

PART 6 – Encryption (Enterprise Security)

Encryption Types:

Encryption	Description
SSE-S3	AWS managed
SSE-KMS	Enterprise-grade

Encryption	Description
Client-side	Application encrypts

Lab 4 – Enable Encryption (KMS)

bash

 Copy code

```
aws s3api put-bucket-encryption \
--bucket company-prod-storage-unique123 \
--server-side-encryption-configuration '{
  "Rules": [
    {
      "ApplyServerSideEncryptionByDefault": {
        "SSEAlgorithm": "AES256"
      }
    }
  ]
}'
```

Enterprise Scenario:

Finance data protection.

PART 7 – Access Control (Enterprise Security Model)

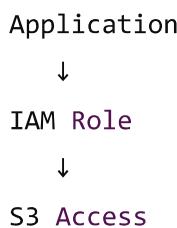
Access Methods:

- IAM Policies
- Bucket Policies
- IAM Roles

Enterprise Model:

pgsql

 Copy code



```

graph TD
    Application[Application] --> IAMRole[IAM Role]
    IAMRole --> S3Access[S3 Access]
  
```

Professional Demo 3 – Bucket Policy Example

Allow EC2 access:

policy.json

json

 Copy code

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Principal": "*",  
      "Action": "s3:GetObject",  
      "Resource": "arn:aws:s3:::company-prod-storage-unique123/*"  
    }  
  ]  
}
```

Apply:

bash

 Copy code

```
aws s3api put-bucket-policy \  
--bucket company-prod-storage-unique123 \  
--policy file://policy.json
```

PART 8 – SDK Integration with S3

Lab 5 – Upload using Python SDK

python

 Copy code

```
import boto3  
  
s3 = boto3.client('s3')
```

```
s3.upload_file(  
    'test.txt',  
    'company-prod-storage-unique123',  
    'logs/test.txt'  
)
```

Enterprise Use Case:

Application uploads logs automatically.

PART 9 – Enterprise Architecture Scenario

Production Application Architecture:

pgsql

 Copy code

Application Server (EC2)

↓

IAM Role

↓

Amazon S3

↓

Lifecycle → Glacier

↓

Backup Storage

PART 10 – Cost Optimization Best Practices

Key Cost Factors:

- Storage class
- Data transfer
- Number of requests

Enterprise Strategy:

Use:

- Intelligent Tiering

- Lifecycle policies
 - Compression
-

Professional Demo 4 – Analyze Storage

bash

 Copy code

```
aws s3 ls s3://company-prod-storage-unique123 --recursive --human-readable --summ
```

End-to-End Enterprise Lab

Participants will:

Step 1

Create S3 bucket

Step 2

Upload files

Step 3

Enable versioning

Step 4

Enable encryption

Step 5

Apply lifecycle policy

Step 6

Upload via SDK

Professional Troubleshooting

Error:

nginx

Copy code

AccessDenied

Fix:

Check IAM permissions.

Error:

nginx

 Copy code

BucketAlreadyExists

Fix:

Use globally unique name.

Enterprise Best Practices

Always:

- Enable versioning
 - Enable encryption
 - Use lifecycle policies
 - Use IAM roles (not keys)
 - Use storage classes wisely
-

Enterprise Architecture Summary

sql

 Copy code

Application

↓

IAM Role

↓

S3 Bucket

↓

Versioning + Encryption



Lifecycle → Archive

Professional Capstone Lab (Enterprise Scenario)

Build enterprise log storage system:

- Secure storage
 - Version protection
 - Cost optimization
 - Automated lifecycle
-

If you want, I can also prepare **Day 5 – ECS Container Deployment with S3 Integration, IAM Roles, Load Balancer, and CI/CD**, which completes real enterprise architecture.



I have DAY 5-> TOC below -> I wanted professional Demos and Lab steps -> also examples, scenarios: Day 5 – Lambda (Part 1) – Functions & Triggers

Lambda Introduction and Use Cases

- What is serverless computing?
 - When to choose Lambda vs ECS/EKS.
- Lambda Execution Model
- Function invocation flow, event sources, cold starts.
 - Memory allocation and timeout configuration.

IAM Role & Permissions

- Execution role vs invocation permissions.
- How to grant Lambda access to Aurora RDS resources.

Networking Basics for Lambda (VPC Access)

- Public vs private Lambda configuration.
- Subnets, Security Groups, and NAT Gateway impact.

Environment Variables & Secrets

- Managing DB credentials via SSM Parameter Store or

Secrets Manager.

- Encrypting environment variables with KMS.

Below is a Day 5 – AWS Lambda (Part 1) Professional Training Plan designed for senior professionals, with enterprise architecture thinking, demos, hands-on labs, production scenarios, security integration, and VPC + RDS connectivity.

This focuses on **real-world enterprise serverless architecture**, not basic hello-world.

Day 5 – AWS Lambda (Part 1) – Functions & Triggers (Enterprise-Level)

Learning Outcomes

By end of Day 5, participants will:

- Understand serverless architecture deeply
 - Know when to use Lambda vs ECS/EKS
 - Create Lambda functions with proper IAM roles
 - Integrate Lambda with S3, API, and RDS
 - Configure Lambda inside VPC securely
 - Manage secrets using Parameter Store / Secrets Manager
 - Implement enterprise-grade security model
-

PART 1 – Serverless Computing (Enterprise Perspective)

What is Serverless?

Serverless means:

You manage code, AWS manages infrastructure.

Traditional model:

mathematica

Copy code

```
Developer  
↓  
Provision EC2  
↓  
Install software  
↓  
Maintain server
```

Serverless model:

css

Copy code

```
Developer  
↓  
Upload code  
↓  
AWS runs automatically
```

Enterprise Architecture Example

Traditional:

pgsql

Copy code

```
Application  
↓  
EC2 Server (always running)
```

Serverless:

vbnet

Copy code

```
Application Event  
↓  
Lambda Function  
↓  
Runs on demand
```

Enterprise Use Cases

Lambda used for:

- File processing
 - API backend
 - Event-driven automation
 - Log processing
 - Database triggers
-

Lambda vs ECS vs EKS (Enterprise Decision Matrix)

Use Case	Lambda	ECS	EKS
Short tasks	Yes	No	No
Long-running services	No	Yes	Yes
Microservices	Yes	Yes	Yes
Containers	No	Yes	Yes
Kubernetes required	No	No	Yes

Professional Demo 1 – Create Lambda Function

Scenario

Application uploads file → Lambda processes file automatically.

Lab 1 – Create Lambda Function

Step 1: Go to AWS Console → Lambda

Click:

Create Function

Choose:

csharp

 Copy code

Author **from scratch**

Enter:

vbnet

 Copy code

Function name: **process-file-function**

Runtime: **Python 3.11**

Choose:

Create new IAM role

Click Create.

PART 2 – Lambda Execution Model

Invocation Flow Architecture

sql

 Copy code

Event Source (S3/API/EventBridge)

↓

Lambda Service

↓

Execution Environment

↓

Function Runs

↓

Returns response

Professional Demo 2 – Lambda Execution Example

Code:

`python` Copy code

```
def lambda_handler(event, context):
    print("Event received:", event)

    return {
        'statusCode': 200,
        'body': 'Lambda executed successfully'
    }
```

Click:

Deploy

Click:

Test

Explain Cold Start (Enterprise Concept)

Cold start occurs when:

Lambda starts new execution environment.

Impacts latency slightly.

Enterprise mitigation:

- Provisioned concurrency
 - Proper memory configuration
-

PART 3 – Memory and Timeout Configuration

Enterprise tuning improves performance.

Example settings:

`makefile` Copy code

Memory: 512 MB

Timeout: 30 seconds

More memory → More CPU → Faster execution

Lab 2 – Configure Memory and Timeout

Go to:

Configuration → General Configuration

Edit:

Memory: 512 MB

Timeout: 30 sec

Save.

PART 4 – IAM Role & Permissions

Lambda uses execution role to access resources.

Architecture:

pgsql

 Copy code

Lambda Function

↓

IAM Execution Role

↓

Access S3 / RDS / DynamoDB

Professional Demo 3 – Add S3 Access Permission

Go to:

Lambda → Configuration → Permissions

Click:

Execution Role → Add policy

Attach:

nginx

 Copy code

AmazonS3FullAccess

PART 5 – Lambda Access to Aurora RDS (Enterprise Scenario)

Architecture:

pgsql

 Copy code

Lambda Function

↓

IAM Role

↓

Aurora RDS

Requirements:

- IAM Role permission
 - VPC access
 - Security group access
-

PART 6 – Lambda Networking (VPC Access)

Enterprise Database is inside private VPC.

Lambda must connect securely.

Architecture:

vbnet

 Copy code

```
Lambda Function
```

```
↓
```

```
Private Subnet
```

```
↓
```

```
Security Group
```

```
↓
```

```
Aurora RDS
```

Lab 3 – Configure Lambda VPC Access

Go to:

Lambda → Configuration → VPC

Select:

- Your VPC
- Private subnet
- Security group

Save.

Enterprise Impact:

Lambda can now access database.

PART 7 – Environment Variables & Secrets

Never hardcode credentials.

Bad:

```
ini
```

 Copy code

```
password = "admin123"
```

Enterprise Solution:

Use:

- Secrets Manager
 - Parameter Store
-

Lab 4 – Add Environment Variable

Go to:

Lambda → Configuration → Environment Variables

Add:

ini

Copy code

```
DB_HOST = database-url
DB_USER = admin
```

Save.

PART 8 – Using Parameter Store (Enterprise Secure Method)

Store secret:

Go to:

SSM → Parameter Store → Create Parameter

Name:

pgsql

Copy code

db-password

Type:

SecureString

Professional Demo 4 – Access Parameter Store via Lambda

Code:

python

 Copy code

```
import boto3

ssm = boto3.client('ssm')

def lambda_handler(event, context):

    response = ssm.get_parameter(
        Name='db-password',
        WithDecryption=True
    )

    password = response['Parameter']['Value']

    print(password)
```

Enterprise Benefit:

Secrets stored securely.

PART 9 – Lambda Trigger Example (S3 Trigger)

Enterprise Scenario:

File upload triggers Lambda.

Architecture:

csharp

 Copy code

User uploads `file`

↓

S3 Bucket

↓

Lambda Triggered

↓

Process `file`

Lab 5 – Create S3 Trigger

Go to:

Lambda → Configuration → Triggers

Add trigger:

Select:

S3

Choose bucket.

Save.

Test by uploading file.

Lambda executes automatically.

PART 10 – Encryption using KMS

Enterprise uses KMS to encrypt secrets.

Architecture:

nginx

 Copy code

Secrets Manager



KMS Encryption



Lambda Access

Professional Demo 5 – Enable KMS Encryption

Go to:

Lambda → Configuration → Environment Variables

Enable encryption using KMS key.

Enterprise End-to-End Architecture Scenario

sql

 Copy code

User uploads file



S3 Bucket



Lambda Function



Fetch secret from Parameter Store



Connect to Aurora RDS



Store processed data

End-to-End Professional Lab

Participants will:

Step 1

Create Lambda function

Step 2

Configure IAM role

Step 3

Configure environment variables

Step 4

Store secrets in Parameter Store

Step 5

Configure VPC access

Step 6

Add S3 trigger

Step 7

Test Lambda execution

Professional Troubleshooting

Error:

nginx

 Copy code

`AccessDeniedException`

Fix:

Add proper IAM permissions.

Error:

nginx

 Copy code

`Timeout error`

Fix:

Increase timeout or fix networking.

Error:

pgsql

 Copy code

Cannot connect to database

Fix:

Check security group rules.

Enterprise Best Practices

Always:

- Use IAM roles
 - Use Parameter Store or Secrets Manager
 - Use private VPC for database access
 - Use KMS encryption
 - Use environment variables
-

Enterprise Architecture Summary

vbnnet

 Copy code

S3 Event

↓

Lambda Function

↓

IAM Role

↓

Secrets Manager

↓

Aurora RDS

Professional Capstone Lab

Build secure serverless file processing system.

If you want, I can prepare Day 6 – Lambda Part 2 with API Gateway integration, real-time enterprise REST API, and production-grade architecture, which is the most important enterprise use case.



I have DAY 6 toc below --> I wanted professional Demos and Lab steps -> also examples, scenarios: -- Day 6 – KMS (Key Management Service) & RDS Integration ()

Why Encryption Matters

- Data at rest vs data in transit.
- Compliance (GDPR, HIPAA, PCI DSS) use cases.
- Shared responsibility for data protection.

Introduction to AWS KMS

- What is AWS Key Management Service?
- Types of Keys:
 - AWS Managed Keys (e.g., aws/rds)
 - Customer Managed Keys (CMK)
 - AWS Owned Keys
- Key components: Key ID, Key Policy, Grants, Aliases, Rotation.

KMS Key Management Concepts

- Envelope encryption concept.
- Key hierarchy (CMK → DEK → Resource Encryption).
- Automatic key rotation.
- Logging with CloudTrail for audit visibility.

Access Control in KMS

- How IAM and Key Policies work together.
- Least privilege for KMS usage (separate encrypt/decrypt roles).

Below is a Day 6 – AWS KMS & RDS Integration (Enterprise-Level Training) designed for senior professionals, with security architecture, encryption design, compliance scenarios, professional demos, hands-on labs, and real-world integration with RDS.

This is critical for enterprise security, compliance, and production-grade encryption architecture.

Day 6 – AWS KMS & RDS Integration (Enterprise-Level)

Learning Outcomes

By end of Day 6, participants will:

- Understand enterprise encryption architecture
 - Understand KMS key types and usage
 - Implement encryption for RDS using KMS
 - Implement secure key policies and IAM access
 - Understand envelope encryption deeply
 - Enable audit logging using CloudTrail
 - Implement least privilege access model
-

PART 1 – Why Encryption Matters (Enterprise Security Foundation)

Enterprise Threat Model

Without encryption:

```
pgsql
```

Copy code

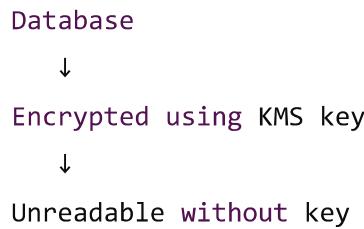
Database → Data stored **in plain text**

If attacker gains access → Data exposed.

With encryption:

```
pgsql
```

Copy code



Data at Rest vs Data in Transit

Type	Description	Example
Data at rest	Stored data encryption	RDS, S3
Data in transit	Network encryption	HTTPS, TLS

Enterprise Example:

pgsql

 Copy code

Application → HTTPS → RDS (encrypted at rest)

Compliance Requirements (Enterprise)

Encryption required for:

- GDPR → User data protection
- HIPAA → Healthcare data
- PCI DSS → Payment data

Enterprise Example:

Payment database MUST use encryption.

PART 2 – AWS Shared Responsibility Model (Encryption Perspective)

AWS responsible for:

- Infrastructure security

Customer responsible for:

- Data encryption
 - Key management
 - Access control
-

PART 3 – AWS KMS Introduction

What is AWS KMS?

AWS KMS is key management service.

Used to:

- Create encryption keys
- Encrypt resources
- Control access to encryption keys

Architecture:

pgsql

 Copy code

Resource (RDS, S3, Lambda)



Data Encryption Key (DEK)



Encrypted using CMK (in KMS)

PART 4 – Types of KMS Keys

1. AWS Managed Keys

Example:

bash

 Copy code

aws/rds

aws/s3

Fully managed by AWS.

Limited control.

2. Customer Managed Keys (CMK) – Enterprise Recommended

Full control:

- Rotation
 - Permissions
 - Policies
-

3. AWS Owned Keys

Used internally by AWS.

Not visible to customer.

Professional Demo 1 – Create Customer Managed Key

Lab 1 – Create CMK

Go to:

AWS Console → KMS → Create Key

Select:

sql

Copy code

Symmetric

Choose:

mathematica

Copy code

Encrypt and decrypt

Enter:

Alias:

vbnets

 Copy code

company-rds-key

Click Create.

PART 5 – Key Components

Component	Purpose
Key ID	Unique identifier
Alias	Human readable name
Key Policy	Defines access
Grants	Temporary access
Rotation	Automatic key change

PART 6 – Envelope Encryption (Critical Enterprise Concept)

Enterprise encryption uses 2-layer encryption.

Architecture:

pgsql

 Copy code

Data

↓

Encrypted using DEK (Data Encryption Key)

↓

DEK encrypted using CMK



Stored securely

Why?

Performance + security.

PART 7 – Key Hierarchy Architecture

mathematica

 Copy code

Customer Managed Key (CMK)



Data Encryption Key (DEK)



Encrypt Resource



RDS Database

PART 8 – Enable Key Rotation

Enterprise security best practice.

Lab 2 – Enable Rotation

Go to:

KMS → Select key → Key rotation

Enable:

mathematica

 Copy code

Automatic rotation

Rotation happens yearly.

PART 9 – Audit Logging using CloudTrail

Enterprise compliance requires audit.

CloudTrail logs:

- Who used key
- When used
- Which resource used

Architecture:

vbnet

Copy code

User

↓

KMS Key Usage

↓

CloudTrail logs event

Professional Demo 2 – View KMS Logs

Go to:

CloudTrail → Event History

Filter:

css

Copy code

Event Source: kms.amazonaws.com

View logs.

PART 10 – Access Control in KMS

Access controlled by:

- IAM policies
- Key policies

Both must allow access.

Architecture:

vbnet

 Copy code

IAM Role

↓

Key Policy

↓

KMS Key

Professional Demo 3 – IAM Policy for KMS

Example Policy:

json

 Copy code

```
{
  "Effect": "Allow",
  "Action": [
    "kms:Encrypt",
    "kms:Decrypt"
  ],
  "Resource": "*"
}
```

PART 11 – RDS Encryption using KMS (Enterprise Scenario)

Enterprise Architecture:

pgsql

 Copy code

Application

↓

RDS Database



Encrypted using KMS CMK

Lab 3 – Create Encrypted RDS Database

Go to:

RDS → Create database

Select:

makefile

 Copy code

Engine: MySQL

Enable:

nginx

 Copy code

Encryption

Choose:

vbnet

 Copy code

company-rds-key

Create database.

Verify Encryption

Go to:

RDS → Database → Configuration

Check:

vbnet

Copy code

Encryption: Enabled
KMS key: company-rds-key

PART 12 – Lambda Access to Encrypted RDS

Enterprise architecture:



Lab 4 – Add Lambda Access to KMS

Attach policy to Lambda role:

```
json
{
  "Effect": "Allow",
  "Action": [
    "kms:Decrypt"
  ],
  "Resource": "*"
}
```

PART 13 – CLI Demo for KMS

Create key using CLI:

bash

 Copy code

```
aws kms create-key \
--description "Enterprise encryption key"
```

List keys:

bash

 Copy code

```
aws kms list-keys
```

PART 14 – SDK Demo using Python

Encrypt data:

python

 Copy code

```
import boto3

kms = boto3.client('kms')

response = kms.encrypt(
    KeyId='alias/company-rds-key',
    Plaintext='Sensitive Data'.encode()
)

print(response['CiphertextBlob'])
```

Enterprise Scenario – Banking Application

Architecture:

pgsql

 Copy code

Banking Application

↓

Lambda

↓

Encrypted RDS



KMS Key



CloudTrail Audit

Compliance satisfied.

End-to-End Enterprise Lab

Participants will:

Step 1

Create CMK

Step 2

Enable key rotation

Step 3

Create encrypted RDS

Step 4

Grant IAM access

Step 5

Test encryption using SDK

Step 6

View audit logs

Professional Troubleshooting

Error:

SCSS

Copy code

AccessDeniedException (KMS)

Fix:

Add IAM permission and key policy.

Error:

```
pgsql|
```

 Copy code

Cannot access encrypted database

Fix:

Grant role access to KMS.

Enterprise Best Practices

Always:

- Use Customer Managed Keys
 - Enable rotation
 - Enable CloudTrail logging
 - Use least privilege access
 - Separate encrypt and decrypt roles
-

Enterprise Architecture Summary

```
mathematica
```

 Copy code

Application

↓

IAM Role

↓

KMS CMK

↓

Encrypt DEK

↓

Encrypt RDS Data

Professional Capstone Lab

Build secure encrypted database architecture using KMS and RDS.

If you want, I can prepare **Day 7 – Real Enterprise Project integrating S3 + Lambda + RDS + KMS + IAM + Terraform**, which simulates actual production system design.

